



Data Modernization

[Home](#)[Services](#)[Analytics Roadmap](#)[Industries](#)[About](#)[Leadership](#)[Contact](#)[Training](#)[Careers](#)[Blog](#)[FAQs](#)

Overview

Durga Analytics helps organizations modernize legacy data platforms into agile, cloud-native ecosystems that drive operational efficiency and real-time decision-making.

Key Capabilities

- Cloud Migration of legacy databases and ETL frameworks
- Data lake & lakehouse architecture setup
- Modern ingestion frameworks (batch & real-time)
- Data cataloging and metadata management
- Modern BI tools and dashboard refresh programs

Approach

- 1. Assessment:** Audit current systems, data sources, usage & quality
- 2. Planning:** Build a cloud migration plan with modernization milestones
- 3. Execution:** Rebuild pipelines, optimize performance, validate outcomes
- 4. Enablement:** Train teams, deploy governance, and automate operations

Benefits

- Accelerated time-to-insight with scalable cloud solutions

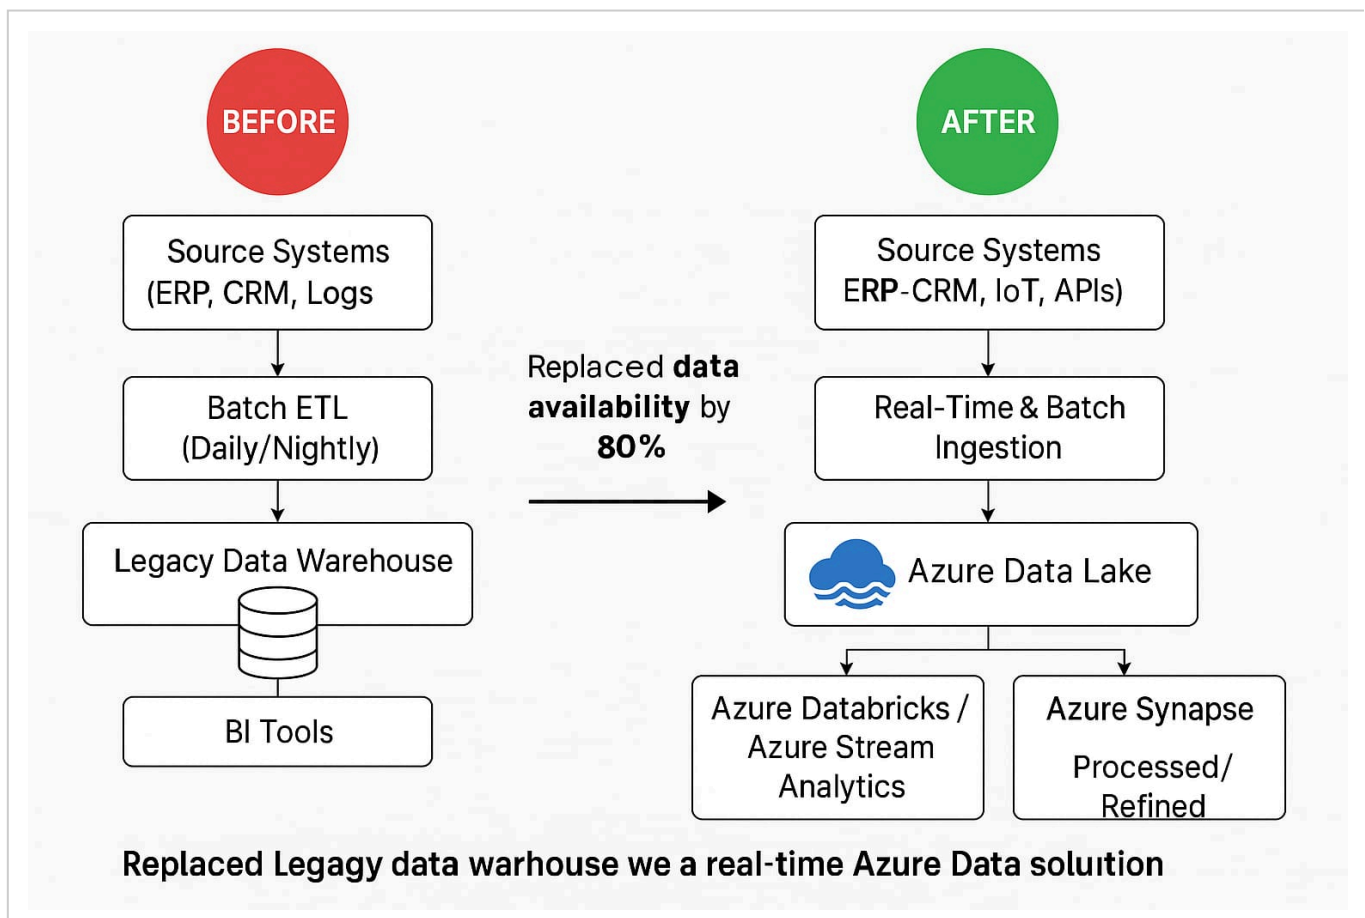
- Reduced total cost of ownership
- Future-proof architecture that supports AI and ML adoption
- Enhanced data accessibility, quality, and governance

Case Highlights

- **Replaced legacy data warehouse with a real-time Azure Data Lake solution, improving data availability by 80%**

Pre Architecture: Monolithic on-premise data warehouse with batch ingestion, limited scalability, and delayed reporting cycles (24–48 hours).

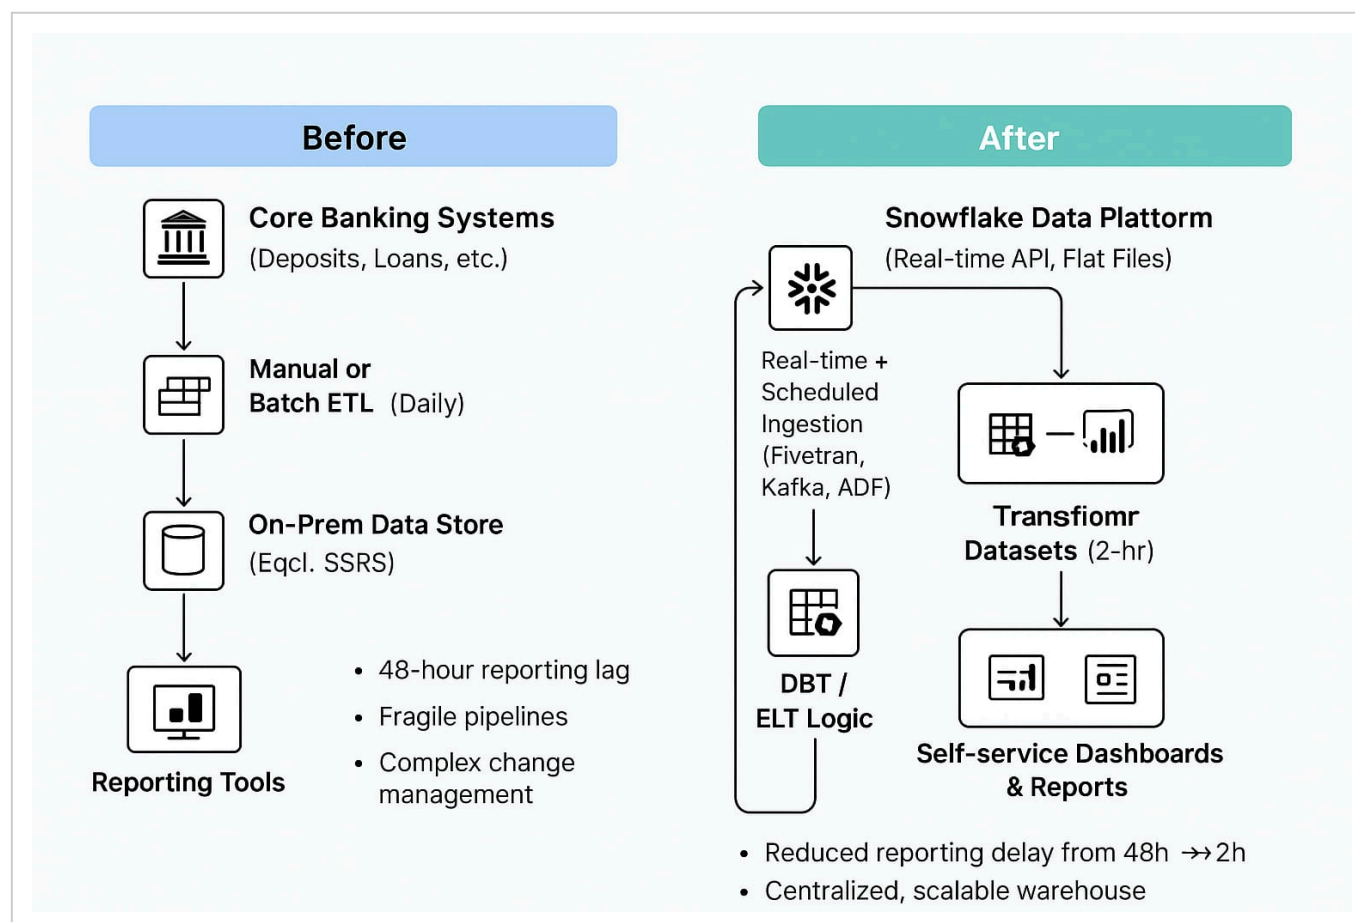
Post Architecture: Cloud-native Azure Data Lake solution with real-time streaming ingestion, scalable compute/storage, and modern analytics stack enabling 80% faster data availability.



- **Built ingestion and reporting platform on Snowflake for a retail bank, reducing reporting delays from 2 days to 2 hours**

Pre Architecture: Fragmented on-prem databases (Core Banking, Cards, CRM) feeding a legacy EDW via nightly/weekly batch ETL (SSIS/hand-coded scripts). Long data latency (T+2 days), duplicate logic in Excel/desktop BI, limited scalability during month-end, manual data quality checks, and slow onboarding of new data sources (4–6 weeks).

Post Architecture: Central Snowflake multi-cluster warehouse on cloud object storage; automated ingestion (change data capture + incremental batch + near real-time streams via Kafka/Fivetran/Snowpipe), standardized raw → curated → semantic layers (dbt / SQL). Orchestrated transformations (Airflow/Prefect) with automated data quality tests (Great Expectations) and role-based security (RBAC + masking). Unified metrics layer powering self-service BI dashboards; SLA-driven pipelines delivering refreshed data every 15–30 minutes (2-hour → near real-time), accelerating analytics & regulatory reporting.



All Services

Strategy & Advisory

- Analytics Roadmap
- Data Strategy
- Platform Strategy

Engineer Your Data

- Data Modernization
- Data Foundation
- Data Operations

Differentiate with AI/ML

- Data Science
- AI Engineering
- ML Products & Platforms

Operationalize Insights

- Experience Consulting
- Application Engineering
- Business Intelligence
- MLOps

Training services

- Databricks Bootcamp
- Snowflake
- ETRM

